

15-213 Recitation 6

# Introduction to Computer Systems

Hasan Al-Jawaheri

2 October, 2013

# Today


- ▶ Memory Hierarchy
- ▶ Caches
- ▶ Exam Review
- ▶ Cachelab

# Caching

- ▶ Temporal Locality (Time)
  - ▶ A memory location accessed is likely to be accessed again multiple times in the future
  - ▶ After accessing address  $X$  in memory, save the bytes in cache for future access
- ▶ Spatial Locality (Space)
  - ▶ If a location is accessed, then nearby locations are likely to be accessed in the future.
  - ▶ After accessing address  $X$ , save the block of memory around  $X$  in cache for future access

# Memory Addresses

- ▶ A *cache* is a set of  $2^s$  cache sets
- ▶ A *cache set* is a set of  $E$  cache lines
  - ▶  $E = 1 \rightarrow$  Direct-mapped
  - ▶  $E$ -way associative
- ▶ Each cache line stores a block
  - ▶ Each block has  $2^b$  bytes



# Miss Rate

- Fraction of memory references not found in the cache

$$\text{Missrate} = 1 - \text{Hitrate}$$

# Real Example!

- ▶ 8 bit-address space
- ▶ Direct-mapped 32-byte cache
  - ▶ 4 byte cache blocks

# Real Example!

- ▶ Indicate which fields do the bits represent
  - ▶ “SI” → Set index
  - ▶ “BO” → Block offset
  - ▶ “CT” → Cache tag



# Real Example!

- ▶ Indicate which fields do the bits represent
  - ▶ “SI” → Set index
  - ▶ “BO” → Block offset
  - ▶ “CT” → Cache tag





# Real Example!

Load No.	Hex Address	Binary Address	Set Number?	Hit or Miss?
1	c7	110 001 11		
2	55	010 101 01		
3	1a	000 110 10		
4	c5	110 001 01		
5	e6	111 001 10		
6	56	010 101 10		
7	77	011 101 11		
8	28	001 010 00		
9	75	011 101 01		
10	94	100 101 00		

# Real Example!

Load No.	Hex Address	Binary Address	Set Number?	Hit or Miss?
1	c7	110 001 11	1	M
2	55	010 101 01	5	M
3	1a	000 110 10	6	M
4	c5	110 001 01	1	H
5	e6	111 001 10	1	M
6	56	010 101 10	5	H
7	77	011 101 11	5	M
8	28	001 010 00	2	M
9	75	011 101 01	5	H
10	94	100 101 00	5	M

# Real Example!

- Which of the following is a possible final state for the cache:

0	1	2	3	4	5	6	7
X	7	1	X	X	4	1	X

0	1	2	3	4	5	6	7
X	1	1	X	0	2	0	X

0	1	2	3	4	5	6	7
X	1	7	X	4	4	0	X

0	1	2	3	4	5	6	7
7	X	1	0	4	4	0	X

0	1	2	3	4	5	6	7
1	7	X	X	1	4	0	X

0	1	2	3	4	5	6	7
X	1	7	X	X	4	4	0

0	1	2	3	4	5	6	7
X	7	1	X	X	4	0	X

# Real Example!

- Which of the following is a possible final state for the cache:

0	1	2	3	4	5	6	7
X	7	1	X	X	4	1	X

0	1	2	3	4	5	6	7
X	1	1	X	0	2	0	X

0	1	2	3	4	5	6	7
X	1	7	X	4	4	0	X

0	1	2	3	4	5	6	7
7	X	1	0	4	4	0	X

0	1	2	3	4	5	6	7
1	7	X	X	1	4	0	X

0	1	2	3	4	5	6	7
X	1	7	X	X	4	4	0

0	1	2	3	4	5	6	7
X	7	1	X	X	4	0	X

# Another Example

This problem requires you to analyze the cache behavior of a function that sums the elements of an array A:

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Assume the following:

- The memory system consists of registers, a single L1 cache, and main memory.
- The **cache is cold** when the function is called and the array has been initialized elsewhere.
- Variables **i, j, and sum are all stored in registers.**
- The array **A is aligned in memory** such that the first two array elements map to the same cache block.
- `sizeof(int) == 4.`
- The cache is **direct mapped**, with a **block size of 8 bytes.**

# Another Example

**CACHE CONSISTS OF 2 SETS**

# Another Example

Address: 00 0 000

Block offset: 000

Set: 0

Tag: 00

1	3	5	7
2	4	6	8

Cold Miss

00	0	000

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Address: 01 0 000

Block offset: 000

Set: 0

Tag: 01

1	3	5	7
2	4	6	8

Miss

01	0	000



# Another Example

Address: 00 0 100

Block offset: 100

Set: 0

Tag: 00

1	3	5	7
2	4	6	8

Miss

00	0	100

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Address: 01 0 100

Block offset: 100

Set: 0

Tag: 01

1	3	5	7
2	4	6	8

Miss

01	0	100

# Another Example

Address: 00 1 000

Block offset: 000

Set: 1

Tag: 00

1	3	5	7
2	4	6	8

Cold Miss

01	0	100
00	1	000

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Address: 01 1 000

Block offset: 000

Set: 1

Tag: 01

1	3	5	7
2	4	6	8

Miss

01	0	100
01	1	000

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Address: 00 1 100

Block offset: 100

Set: 1

Tag: 00

1	3	5	7
2	4	6	8

Miss

01	0	100
00	1	100

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Address: 01 1 100

Block offset: 100

Set: 1

Tag: 01

1	3	5	7
2	4	6	8

Miss

01	0	100
01	1	100

# Another Example

```
int A[2][4];  
int sum()  
{  
    int i, j, sum=0;  
    for (j=0; j<4; j++) {  
        for (i=0; i<2; i++) {  
            sum += A[i][j];  
        }  
    }  
    return sum;  
}
```

Everything missed!

# Another Example

**CACHE CONSISTS OF 4 SETS**



# Another Example

## ► Problem in previous cache:

- Every time  $i$  changes, the new address has the same set but a different tag
- Then we will miss every time

1	3	5	7
2	4	6	8

## ► New cache:

- Every time  $i$  changes, the new address has a different set
- But notice, for 3, 4, 7 and 8 we won't miss because 1, 2, 5 and 6 (respectively) loaded them.
- **Miss rate is only half!**

1	3	5	7
2	4	6	8

# Cachelab

- ▶ Due: Sometime after Eid \o/
- ▶ Office Hours/Piazza

# Final

- ▶ Integers
  - ▶ Edge cases
  - ▶ Casting
  - ▶ Endianness
- ▶ Floating Points
- ▶ Assembly
  - ▶ How loops work
  - ▶ How addressing and dereferencing work
  - ▶ Typical address calculation scheme
  - ▶ How to use the book for quick function lookup

- ▶ Structs and unions
- ▶ Stack
  - ▶ Frames
  - ▶ Stack management
  - ▶ Caller/Callee roles
- ▶ Caches

FOCUS ON THE DETAILS!